

# Modern VLSI Design — Chapter 1: Overview, Trends, and Design Flow

EECS 168 • Big-picture context, abstraction ladder, IP methodology, and scaling realities.

## 0) Map of the terrain (what Chapter 1 is really doing)

- **Why VLSI:** Integration brings speed, power, and cost wins at product scale.
- **The scaling story:** Moore's Law history, the end of classic Dennard scaling, and the *power wall*.
- **Design at scale:** Abstraction levels, domains, and the modern VLSI flow.
- **IP as leverage:** Hard vs. soft IP, lifecycle, and practical selection/qualification.
- **Reliability & reality:** Nanometer reliability issues and design-for-yield/manufacturing themes.

### Key Idea

VLSI is a two-front war: *technology scaling* (physics, process) and *design scaling* (abstraction, hierarchy, IP). You need both to ship chips on time.

## 1) Why VLSI matters (integration economics & applications)

Integration wins:

- **Speed:** lower on-chip interconnect parasitics vs. board-level wiring.
- **Energy:** shorter wires + fewer I/O transitions reduce power.
- **Size & cost:** fewer parts, higher reliability, better manufacturability.

**Where VLSI shows up:** CPUs/MCUs, memory (SRAM/DRAM), I/O/PHY, and accelerators (DSP, vision, ML). Systems are full SoCs, not just processors. *Custom silicon* is widely used across phones, datacenter, networking, and edge devices. [Slides: overview of applications and industry trend] (see sources at end)

## 2) A quick historical arc (just enough context)

From early mechanical computation → vacuum tubes → the transistor (1947–48) → integrated circuits → micro-processors (Intel 4004, 1971) → deep sub-micron SoCs. Frequency ramped hard through the 90s/early 2000s before thermal and energy walls redirected progress toward parallelism and specialization. [Slides: history snapshots, frequency/power charts]

### Gotcha

Raw frequency scaling lost steam; the modern playbook is *parallelism* (multicore), *specialization* (accelerators), and *system design* (memory, packaging, interconnects).

### 3) Scaling realities: Moore, Dennard, and the power wall

#### 3.1 Moore's Law (transistor count)

For decades, transistors/chip doubled every  $\sim 18\text{--}24$  months, enabling more cores, cache, and features. [Slide index shows Moore's Law content]

#### 3.2 Frequency and power density

Lead microprocessor **frequency** rose rapidly, but **power** and **power density** rose too; delivering and dissipating that power became prohibitive. [Frequency/Power/Power-density charts]

##### Key Idea

**Power wall:** Frequency stalled around a few GHz due to thermal limits; pushing V or f further exploded power density, overwhelming cooling budgets.

#### 3.3 Dennard scaling ended

Classic Dennard: as devices shrink, power density stays constant (voltage scales with length). That assumption broke; leakage rose, voltage scaling slowed, power density soared. [Slides: "End of Moore/Dennard/Power Wall"]

##### Checklist

###### Implications:

- Energy efficiency is king: performance/W matters more than raw f.
- Architecture must carry more: multicore, data-parallel units, accelerators.
- Packaging/interconnect innovations (chiplets, advanced 2.5D/3D) join the critical path.

##### Gotcha

"Just crank the clock" is done. Gains now come from *doing less work* per result (specialization) or *doing work in parallel* at lower V/f.

### 4) The VLSI design problem (why we need abstraction)

**Conflicting objectives:** area, performance, energy/power, reliability, testability, and time-to-market clash. Designs span *multiple abstraction levels*, from transistors up to CPUs. [Slides: "Challenges in VLSI design"]

##### Key Idea

**Divide-and-conquer:** limit what a human must reason about at once. Use hierarchy (compose from verified blocks) and abstraction (only the right detail per level).

### 5) Abstraction levels and design domains (the ladder)

#### 5.1 Abstraction ladder

- **Specification:** English/system models, performance and cost goals.

- **Behavioral:** executable spec/high-level synthesis targets.
- **Register-Transfer (RTL):** state machines, datapaths, cycle-level timing.
- **Logic:** gate-level structures, Boolean transforms, basic timing arcs.
- **Circuit:** transistor-level choices, device sizing, analog effects.
- **Layout:** geometric shapes, layers, design rules, exact parasitics.

*Each level has its own modeling objects and metrics (e.g., throughput vs. nanoseconds vs. microns).*

[Slides: design abstractions levels]

## 5.2 Domains & methodologies

Top-down refines specs to masks; bottom-up builds reusable components. Real flows interleave both (e.g., pick a verified memory macro, then architect around it).

[Slides: design domains]

(Top-down/Bottom-up)]

### Checklist

**At each level know:** key objects (e.g., FSM, netlist, layout), validation method (simulation, STA, DRC/LVS), and handoff artifacts (constraints, libs, LEF/DEF, GDS).

## 6) The end-to-end VLSI flow (what happens when)

1. **Specification:** function, PPA targets, interfaces, cost constraints.
2. **Architecture:** partition into blocks; choose memory hierarchies, datapaths; estimate PPA.
3. **Logic/RTL:** write/verify RTL, define protocols, pipeline for timing.
4. **Circuit:** size devices, select standard cells/Vt options; clock tree and IO planning begin.
5. **Physical (layout):** floorplan, place, route, extract parasitics, close timing/power/SI.

**Design validation** happens continuously: forward check between levels; back-annotate real timing/power upstream once parasitics land. **Manufacturing test** (after fab) screens for defects and speed-grades parts; it is *not* the same as design validation.

[Slides: design process; design validation;

manufacturing test]

### Key Idea

Fixing a bug gets exponentially more expensive the later you find it. Validate early, validate often.

## 7) IP-based design (how big chips ship on time)

### 7.1 Types of IP

**Hard IP:** finished layout (fixed pins/layers); process-qualified; great for tight PPA and analog/PHYs. Must meet layout pin placement, layer usage, and transistor sizing standards.

**Soft IP:** RTL/gate netlist; you synthesize/P&R; timing is only firm after layout; wrap to meet interface standards.

[Slides: IP, hard vs. soft IP]

## 7.2 IP across the hierarchy and the lifecycle

**Across hierarchy:** standard cells, memories, RTL blocks, CPUs, buses, I/O devices.

**Lifecycle:** spec → HDL design → validation & documentation/extraction → qualification → IP database/modules → chip integration.

**Using IP:** identify candidates (vendor, open, internal), evaluate suitability, license as needed, and consider qualification if pushing analog limits. [Slides: IP across hierarchy; IP lifecycle; Using IP]

### Checklist

When choosing IP: hard/soft, functionality, PVT corners, power, DRC/LVS readiness, interface compliance, deliverables (models, libs, views), support policy.

## 8) Reliability and yield (design for the real world)

Nanometer techs bring EM (electromigration), SM, NBTI, hot carriers, TDDDB concerns. DFM/DFY practices and architecture/circuit techniques can compensate for weaker devices or variability to improve yield. [Slides: Reliability; DFM/DFY]

### Gotcha

Running closer to physical limits without reliability guardrails saves area today and costs you *lifetime* and *RMA* tomorrow.

## 9) Managing complexity (how we keep our sanity)

**Divide-and-conquer:** group small components into larger modules; reuse proven blocks; maintain clean interfaces. [Slides: Dealing with complexity]

**Hierarchical names & netlists:** instances are typed (e.g., many adders of one type). Hierarchical names (`top/u_core/u_alu.sum`) disambiguate signals across levels. Netlists connect pins to nets; component lists bind instance types. [Slides: hierarchical names, component hierarchy, netlists]

### Key Idea

Good hierarchy and naming aren't bureaucracy—they're how timing closure and debugging remain tractable on million-gate designs.

## 10) Power: the quick working model (for context)

**Dynamic**  $P_{\text{dyn}} \approx \alpha C V^2 f$  (activity  $\alpha$ , switched capacitance  $C$ , supply  $V$ , clock  $f$ ).

**Static**  $P_{\text{leak}} \approx I_{\text{leak}} V$  (dominant in advanced nodes).

**Why it matters here:** Chapter 1's frequency/power/power-density charts motivate all later emphasis on low-power architecture and physical design choices.

### Checklist

**Levers:** reduce  $V$ , reduce  $C$ , reduce  $f$ , reduce activity  $\alpha$ , power-gate blocks, use higher-Vt cells where timing allows, and move compute where energy is cheapest.

## 11) What to be able to do after this chapter

- Walk the abstraction ladder and name artifacts/metrics for each level.
- Explain Moore vs. Dennard vs. the power wall in one minute.
- Describe the VLSI flow end-to-end and where validation vs. manufacturing test fits.
- Distinguish hard vs. soft IP; outline an IP evaluation checklist.
- Name the major reliability risks at nanometer nodes and why DFY/DFM exists.

## Appendix: Mini-Worksheet (practice, keep unsolved)

### A) Abstractions drill

For each level (spec, behavioral, RTL, logic, circuit, layout):

(i) name the primary **modeling object**, (ii) one **metric**, (iii) one **validation method**, (iv) the **handoff artifact** to the next stage.

### B) IP decision exercise

You need a PCIe Gen4 x8 controller. Draft a 7-point checklist comparing a vendor hard IP vs. an open soft IP for your node and PPA targets. Include: deliverables, PVT coverage, interface wrappers, timing views, qual data, support, and licensing.

### C) Scaling story in 6 sentences

Write six sentences that connect: Moore's Law → frequency rise → power density → power wall → multicore/accelerators → packaging/interconnect innovation.

### D) Reliability scan

Pick two risks (e.g., EM, NBTI). State what physical stress causes them and one circuit/architecture mitigation for each.